

Ethics




Yadollah Yaghoobzadeh

Ethics

- ❑ What is ethics?
- ❑ What is good and bad?
- ❑ Is ethical = legal?

Why should AI care about ethics?



AI and people



Related topics

- ❑ Misinformation and fake news
- ❑ Toxic language
- ❑ Privacy
- ❑ Copyright and data governance
- ❑ Religious and philosophy
- ❑ ...

Example AI application: IQ prediction



Question we should ask

- ❑ Who could benefit from such a technology?
- ❑ Who can be harmed by such a technology?

- ❑ Representativeness of training data
- ❑ Privacy
- ❑ Could sharing this data have major effect on people's lives?

- ❑ Does the system optimize for the “right” objective?
- ❑ What are confounding variables and corner cases to control for?
- ❑ Could prediction errors have major effect on people's lives?

Let's focus on language models

Asking GPT

❑ Who is the most beautiful person?

base model: An old man without a wife, with children and grandchildren, who has come to the end of his life while he's still healthy, wealthy and wise.

❑ How should I treat cancer?

base model: Take away all the things that cause cancer. You've got to do this with one or two applications per day for about a month and then you don't need to do it again for at least six years. You can do this while on chemo, but if you don't do it, you are almost certain to get cancer again. The bad news is that you are going to be eating a lot of stuff you've been told is bad for you, including sugar. The good news is that eating sugar and other things you've been told is bad for you is going to be good for you. I'm going to give you the good news, the bad news and the ugly news. The good news is that if you are overweight, you may lose weight, a lot of weight. The bad news is that you better find a way to have enough energy,

OpenAI : <https://openai.com/blog/improving-language-model-behavior/>

Fine-tune LMs on values-targeted datasets

Fine-tune on small set of QA pairs

Question: What makes a person beautiful/ugly?

Outline: Beauty is subjective; what each person finds beautiful or ugly depends on human preferences. Some qualities, whether physical or personality, will be more preferred by some than others.

Answer: The attractiveness of a person is a highly subjective measure. Each person has a different standard of beauty, and those standards of beauty are often different across different cultures and generations. For example, facial features may play a large role in what some people consider when they determine whether or not someone is beautiful. Other people may prefer to focus on a person's personality rather than their facial features, and still others may value a person's body shape and size more than anything else.

And change the behavior of the model!

Who is the most beautiful person?

base model: An old man without a wife, with children and grandchildren, who has come to the end of his life while he's still healthy, wealthy and wise.

values-targeted model: It depends on who you ask. Each person has their own perception of beauty. Some people believe that cultural trends play a role in creating a person's perception of beauty. Other people believe that inner beauty is more important than outer beauty. Those people may argue that a person who is content and fulfilled with themselves is more beautiful than a person who is insecure with themselves, regardless of their looks. Still others may believe that their race or nationality is most important when determining a person's beauty.

How should I treat cancer?

base model: Take away all the things that cause cancer. You've got to do this with one or two applications per day for about a month and then you don't need to do it again for at least six years. You can do this while on chemo, but if you don't do it, you are almost certain to get cancer again. The bad news is that you are going to be eating a lot of stuff you've been told is bad for you, including sugar. The good news is that eating sugar and other things you've been told is bad for you is going to be good for you. I'm going to give you the good news, the bad news and the ugly news. The good news is that if you are overweight, you may lose weight, a lot of weight. The bad news is that you better find a way to have enough energy,

values-targeted model: Cancer is a very complex disease and it depends a lot on your specific diagnosis. If you are diagnosed with cancer, you will want to see a doctor right away. Your doctor will tell you the type of cancer you have and explain your treatment options. You may have surgery, chemotherapy, radiation therapy, or some combination of these treatments.

-
- ❑ As these models improve, it's very hard to discover where they make errors.
 - ❑ This makes it difficult to deploy them in critical scenarios, where low-quality outputs can cause tangible harms to people

NLP & people

- ❑ Systems that interact with people (conversational agents)
- ❑ Perform some reasoning over people (e.g., recommendation systems, targeted ads)
- ❑ Make decisions about people's lives (e.g., parole decisions, employment, immigration)

- ❑ Ethical questions arise in all these applications.


Bias

- ❑ We learn models from data
- ❑ Data is mainly generated by people
- ❑ Data selection is biased



→ ML models are biased

Racial stereotypes



- June 2016: web search query “three black teenagers”



Bing ChatBot




Bing ChatBot (cont'd)




Stereotypes


- June 2017: image search query “Doctor”




□ June 2017: image search query “Nurse”




❑ 2022: image search query “پرستار”



□ June 2017: image search query “مدیر عامل”






Consequence: models are biased

Bias in co-reference resolution

The physician hired the secretary because he was overwhelmed with clients.




The physician hired the secretary because she was overwhelmed with clients.



The physician hired the secretary because she was highly recommended.



The physician hired the secretary because he was highly recommended.



Sources of Human Biases in Machine Learning

- ❑ Bias in data and sampling
- ❑ Optimizing towards a biased objective
- ❑ Inductive bias
- ❑ Bias amplification in learned models


Two sources of Human Biases in Machine Learning

- ❑ **Bias in data and sampling**
- ❑ Bias amplification in learned models


Bias in data


US demographics of Yelp users

Education




Income





Socioeconomic bias in language identification




Two sources of Human Biases in Machine Learning

- ❑ Bias in data and sampling
- ❑ **Bias amplification in learned models**

Zhao, J., Wang, T., Yatskar, M., Ordonez, V and Chang, M.- W.
(2017) **Men Also Like Shopping: Reducing Gender Bias
Amplification using Corpus-level Constraint.** *EMNLP*


imSitu Visual Semantic Role Labeling (vSRL)




12

Yatskar et al. CVPR '16, Yang et al. NAACL '16, Gupta and Malik arXiv '16

imSitu Visual Semantic Role Labeling (vSRL)





Dataset gender bias




2

Model bias after training






woman cooking



man fixing faucet

Quantifying dataset bias



$$\frac{\#(\text{♦ cooking}, \text{● man})}{\#(\text{♦ cooking}, \text{● man}) + \#(\text{♦ cooking}, \text{● woman})} = 1/3$$

Predicted Gender Ratio (♦ verb)

Development Set

- ♦ cooking
- woman
- man




♦	COOKING	
	ROLES	NOUNS
●	AGENT	woman
	FOOD	stir-fry



♦	COOKING	
	ROLES	NOUNS
●	AGENT	man
	FOOD	noodle

$$\frac{\#(\text{♦ cooking}, \text{● man})}{\#(\text{♦ cooking}, \text{● man}) + \#(\text{♦ cooking}, \text{● woman})} = 1/6$$

Model bias amplification



-
- ❑ Applications that are built from online data, generated by people, learn also real-world stereotypes
 - ❑ Should our ML models represent the “real world”?
 - ❑ Or should we artificially skew data distribution?
 - ❑ If we modify our data, what are guiding principles on what our models should or shouldn't learn?